

CIS8695 Final Project

Predicting Breast Cancer malign or Benign using different predicting models

**Submitted By –
Shekha Saxena**

Table of Contents

1. Project Description	3
2. Data Exploration and Preprocessing	4
3. Model – Logistic	6
4. Model – Random Forest	6
5. Model – KNN	9
6. Model – Neural Network	11
7. Results and Discussions	12
8. Summary	13
9. References	13

Project Description

Breast cancer (BC) is one of the most common cancers among women worldwide. The early diagnosis of BC can improve the prognosis and chance of survival significantly, as it can promote timely clinical treatment to patients. Further accurate classification of benign tumours can prevent patients undergoing unnecessary treatments. Thus, the correct diagnosis of BC and classification of patients into malignant or benign groups is the subject of much research. Because of its unique advantages in critical features detection from complex BC datasets, machine learning (ML) is widely recognized as the methodology of choice in BC pattern classification and forecast modelling. We developed 4 robust classifiers to predict if a tumor is malign or benign based on the dataset recorded by University of Wisconsin Hospital at Madison which includes information regarding the texture, size etc. of the tumor. In this project, we reviewed the traditional and advanced deep learning approaches, and recent advances in this field. We show that with sufficient preprocessing and selecting the right feature set, and resampling to address imbalanced datasets, even simple logistic regression and neural network models give superior performance.

Dataset and features

The dataset has the following features represented in columns:

Id - ID number

Diagnosis - The diagnosis of breast tissues (M = malignant, B = benign)

radius_mean - mean of distances from center to points on the perimeter

texture_mean - standard deviation of gray-scale values

perimeter_mean - mean size of the core tumor

area_mean

smoothness_mean - mean of local variation in radius lengths

compactness_mean - mean of $\text{perimeter}^2 / \text{area} - 1.0$

concavity_mean - mean of severity of concave portions of the contour

concave points_mean - mean for number of concave portions of the contour

symmetry_mean

fractal_dimension_mean - mean for "coastline approximation" - 1

radius_se - standard error for the mean of distances from center to points on the perimeter

texture_se - standard error for standard deviation of gray-scale values

perimeter_se

area_se

smoothness_se - standard error for local variation in radius lengths

compactness_se - standard error for $\text{perimeter}^2 / \text{area} - 1.0$

concavity_se - standard error for severity of concave portions of the contour

concave points_se - standard error for number of concave portions of the contour

symmetry_se

fractal_dimension_se - standard error for "coastline approximation" - 1

radius_worst - "worst" or largest mean value for mean of distances from center to points on the perimeter

texture_worst - "worst" or largest mean value for standard deviation of gray-scale values

perimeter_worst

area_worst

smoothness_worst - worst" or largest mean value for local variation in radius lengths

compactness_worst - "worst" or largest mean value for $\text{perimeter}^2 / \text{area} - 1.0$

concavity_worst - "worst" or largest mean value for severity of concave portions of the contour

concave points_worst - "worst" or largest mean value for number of concave portions of the contour

symmetry_worst

fractal_dimension_worst - "worst" or largest mean value for "coastline approximation" – 1

Data Exploration and Preprocessing

For data preprocessing, firstly we need to check if there are any null values in the dataset, which we do in R as below:

```
#check missing values  
colSums(is.na(mydata))
```

From the above result we get to know that, while fetching the dataset there is a column X which is also being fetched, and we do not require it. There is another column ID – which we do not require as it's the ID number of the patient and will not help in predicting if a cancer is malign or benign.

```
#removing unnecessary data columns  
mydata <- mydata[, -c(1, 33)]
```

On checking values of the dataset using `head(mydata)` we find that the values of the category diagnosis is subjective i.e. it's "M" – for Malign and "B" – for Benign. For this reason we need to change the values for Malign to be as 1 and Benign to be as 0.

We do this because we need our data to be categorical in order to use categorical predictive models such as Logistic, etc. We make this change in the dataset with the following command

```
#changing to binary  
mydata$diagnosis = ifelse( mydata$diagnosis=="M", 1, 0)
```

Now, we need to check multicollinearity between the features. We need to check and remove variable showing multicollinearity as they will not help us in making and predicting a robust model. We can check multicollinearity by plotting a graph and by using a function in R which are as follows:

```
#check multicollinearity
ggcorr(mydata[, -1],
       label = TRUE,
       label_alpha = TRUE)
```

```
cor(mydata[, -1])
```


As we can see there are many columns interdependent on each other, hence, showing multicollinearity.

Looking at the matrix, we can immediately verify the presence of multicollinearity between some of our variables. For instance, the **radius_mean** column has a correlation of 1 and 0.99 with **perimeter_mean** and **area_mean** columns, respectively. This is probably because the three columns essentially contain the same information, which is the physical size of the observation (the cell). Therefore, we should only pick one of the three columns when we go into further analysis.

Another place where multicollinearity is apparent is between the "mean" columns and the "worst" column. For instance, the **radius_mean** column has a correlation of 0.97 with the **radius_worst** column. In fact, each of the 10 key attributes display very high (from 0.7 up to 0.97) correlations between its "mean" and "worst" columns. This is somewhat inevitable,

because the "worst" columns are essentially just a subset of the "mean" columns; the "worst" columns are also the "mean" of some values (the three largest values among all observations). Therefore, I think we should discard the "worst" columns from our analysis and only focus on the "mean" columns.

Similarly, it seems like there is multicollinearity between the



attribute's **compactness**, **concavity**, and **concave points**. Just like what we did with the size attributes, we should pick only one of these three attributes that contain information on the shape of the cell. I think **compactness** is an attribute name that is straightforward, so I will remove the other two attributes.

We will now go head and drop all unnecessary columns.


```
#removing all worst columns, perimeter and area
```

```
mydata <- mydata[ , -c(4,5,14,15,22:31)]
```

```
#removing concavity and concave points columns
```

```
mydata <- mydata[ , -c(6,7,14,15)]
```

After removing all these columns we still see that the correlation between compactness_se and fractal_dimension_se is above 80% and since we know that a correlation above 80% is to be removed, we still try to create a model using this, and if we find it of no use, we can remove it later.



For all the models we will be setting our training dataset with 60% of the original data and the validation dataset with the 40% of the original data. We do so in order make our model familiar with the patterns in the dataset with the 60% of the training data.

```
#plot diagnosis count
table(mydata$diagnosis)
barplot(table(mydata$diagnosis), main="Diagnosis",
          xlab="0 - Benign, 1 - Malign", density= 569)
```

```
> table(mydata$diagnosis)
 0    1 
357 212 
> |
```

Models

Logistic Regression:

```
logit.reg <- glm(diagnosis ~ ., data = train, family = "binomial")
```

```
options(scipen=999) # remove scientific notation
summary(logit.reg)
```

```
logit.reg.pred <- predict(logit.reg, valid, type =
"response")
library(caret)
confusionMatrix(as.factor(ifelse(logit.reg.pred > 0.5,
1, 0)), as.factor(valid$diagnosis))
```

Random Forest:

```
rf <- randomForest(as.factor(diagnosis) ~ ., data =
train, ntree = 500,
                  mtry = 4, nodesize = 5, importance =
TRUE)
summary(rf)

summary(tree)
```

Confusion Matrix and Statistics

```

              Reference
Prediction   0    1
           0 137    7
           1   8   76

      Accuracy : 0.9342
      95% CI   : (0.8938, 0.9627)
No Information Rate : 0.636
P-Value [Acc > NIR] : <0.0000000000000002

      Kappa : 0.8583
McNemar's Test P-Value : 1

Sensitivity : 0.9448
Specificity : 0.9157
Pos Pred Value : 0.9514
Neg Pred Value : 0.9048
Prevalence : 0.6360
Detection Rate : 0.6009
Detection Prevalence : 0.6316
Balanced Accuracy : 0.9302

'Positive' Class : 0
```

```
> summary(tree)
left daughter  right daughter  split var  split point  status
Min.   : 0.000  Min.   : 0.00  compactness_se  : 3  Min.   : 0.0000  Min.   : -1.00000
1st Qu.: 0.000  1st Qu.: 0.00  radius_mean    : 3  1st Qu.: 0.0000  1st Qu.: -1.00000
Median : 0.000  Median : 0.00  radius_se     : 3  Median : 0.0000  Median : -1.00000
Mean   : 9.243  Mean   : 9.73  texture_mean   : 3  Mean   : 2.8273  Mean   : -0.02703
3rd Qu.:18.000  3rd Qu.:19.00  fractal_dimension_mean: 2  3rd Qu.: 0.2467  3rd Qu.: 1.00000
Max.   :36.000  Max.   :37.00  (Other)       : 4  Max.   :24.4700  Max.   : 1.00000
NA's                      :19

prediction
Length:37
Class :character
Mode :character
```


```
head(rf$votes,10)
```

```
> head(rf$votes,10)
      0      1
152 0.8950276 0.104972376
212 0.9942857 0.005714286
325 0.9948454 0.005154639
515 0.8351064 0.164893617
114 0.8564103 0.143589744
507 0.6404494 0.359550562
532 0.9455782 0.054421769
372 0.7572816 0.242718447
353 0.0000000 1.000000000
35  0.1164021 0.883597884
> |
```

```
## Plot forest by prediction errors
```


```
plot(rf, type = "simple")
```

```
> summary(rf)
      Length Class  Mode
call           7  -none- call
type           1  -none- character
predicted      341  factor numeric
err.rate      1500  -none- numeric
confusion       6  -none- numeric
votes         682  matrix numeric
oob.times      341  -none- numeric
classes        2  -none- character
importance      48  -none- numeric
importanceSD    36  -none- numeric
localImportance 0  -none- NULL
proximity       0  -none- NULL
ntree           1  -none- numeric
mtry            1  -none- numeric
forest         14  -none- list
y              341  factor numeric
test           0  -none- NULL
inbag           0  -none- NULL
terms           3  terms  call
> |
```

```
## variable importance plot
varImpPlot(rf, type = 1)
```

```
## confusion matrix
rf.pred <- predict(rf, valid)
library(caret)
confusionMatrix(rf.pred, as.factor(valid$diagnosis))
```



Confusion Matrix and Statistics

```

              Reference
Prediction  0   1
0      132  14
1       7   75

      Accuracy : 0.9079
      95% CI   : (0.8627, 0.9421)
No Information Rate : 0.6096
P-Value [Acc > NIR] : <0.0000000000000002

      Kappa : 0.8037
McNemar's Test P-Value : 0.1904

      Sensitivity : 0.9496
      Specificity : 0.8427
      Pos Pred Value : 0.9041
      Neg Pred Value : 0.9146
      Prevalence : 0.6096
      Detection Rate : 0.5789
      Detection Prevalence : 0.6404
      Balanced Accuracy : 0.8962

      'Positive' Class : 0

```

KNN:

Find optimal K

```

set.seed(502)
grid1 <- expand.grid(.k = seq(2, 20, by = 1))
control <- trainControl(method = "cv")
knn.train <- train(diagnosis ~ ., data = train,
                  method = "knn",
                  trControl = control,
                  tuneGrid = grid1)
knn.train

knn.pred <- predict(knn.train, newdata = valid)
confusionMatrix(factor(knn.pred, levels = 0:1),
factor(valid$diagnosis, levels = 0:1))

```

k-Nearest Neighbors

341 samples
12 predictor predict(object, ...)

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 306, 307, 307, 307, 307, 307, ...
Resampling results across tuning parameters:

k	RMSE	Rsquared	MAE
2	0.3226269	0.5721220	0.1481513
3	0.3036478	0.6096939	0.1466106
4	0.3069032	0.5987462	0.1561134
5	0.3083471	0.5977289	0.1612773
6	0.2986795	0.6159746	0.1597759
7	0.2935133	0.6268307	0.1591717
8	0.2880171	0.6400114	0.1605462
9	0.2796651	0.6584481	0.1564052
10	0.2826721	0.6511814	0.1583529
11	0.2815875	0.6545363	0.1598803
12	0.2810938	0.6575058	0.1600210
13	0.2780532	0.6647055	0.1589916
14	0.2790210	0.6628088	0.1606122
15	0.2779876	0.6660904	0.1610476
16	0.2761792	0.6694314	0.1606985
17	0.2771487	0.6682215	0.1626149
18	0.2772171	0.6689187	0.1640110
19	0.2775585	0.6677687	0.1648253
20	0.2773899	0.6677980	0.1658613

RMSE was used to select the optimal model using the smallest value.

Confusion Matrix and Statistics


Reference
Prediction 0 1
0 62 3
1 0 40

Accuracy : 0.9714
95% CI : (0.9188, 0.9941)
No Information Rate : 0.5905
P-Value [Acc > NIR] : <0.0000000000000002

Kappa : 0.9403
McNemar's Test P-Value : 0.2482

Sensitivity : 1.0000
Specificity : 0.9302
Pos Pred Value : 0.9538
Neg Pred Value : 1.0000
Prevalence : 0.5905
Detection Rate : 0.5905
Detection Prevalence : 0.6190
Balanced Accuracy : 0.9651

'Positive' Class : 0



```
# Different distance weighting
#install.packages("kknn")
library(kknn)
set.seed(123)
kknn.train <- train.kknn(diagnosis ~ ., data = train, kmax = 25,
  distance = 2,
  kernel = c("rectangular", "triangular", "epanechnikov"))
plot(kknn.train)
```

kknn.train

```
> kknn.train

Call:
train.kknn(formula = diagnosis ~ ., data = train, kmax = 25, distance = 2, kernel = c("rectangular", "triangular", "epanechnikov"))

Type of response variable: continuous
minimal mean absolute error: 0.08211144
Minimal mean squared error: 0.04709089
Best kernel: triangular
Best k: 5
> |
```

```
kknn.pred <- predict(kknn.train, newdata = valid)
confusionMatrix(factor(kknn.pred, levels = 0:1), factor(valid$diagnosis, levels = 0:1))
```

Confusion Matrix and Statistics

```

      Reference
Prediction 0  1
0    107   2
1     1  56

```

```

      Accuracy : 0.9819
      95% CI : (0.9481, 0.9963)
    No Information Rate : 0.6506
    P-Value [Acc > NIR] : <0.0000000000000002

```

```

      Kappa : 0.9601
    McNemar's Test P-Value : 1

```

```

      Sensitivity : 0.9907
      Specificity : 0.9655
    Pos Pred Value : 0.9817
    Neg Pred Value : 0.9825
      Prevalence : 0.6506
    Detection Rate : 0.6446
    Detection Prevalence : 0.6566
    Balanced Accuracy : 0.9781

```

```

'Positive' Class : 0

```

> |

Neural Network:

```
nn <- neuralnet(diagnosis ~ ., data = train, hidden = c(4,2), linear.output = FALSE)
plot(nn)
```

```

preds.valid <- compute(nn, valid[, -c(1)])
preds.valid.class <- ifelse(preds.valid$net.result > 0.5, 1, 0)
confusionMatrix(as.factor(preds.valid.class), as.factor(valid$diagnosis))

```

Confusion Matrix and Statistics

```

      Reference
Prediction 0  1
0    131  10
1     8  79

```

```

      Accuracy : 0.9211
      95% CI : (0.8781, 0.9525)
    No Information Rate : 0.6096
    P-Value [Acc > NIR] : <0.0000000000000002

```

```

      Kappa : 0.8335
    McNemar's Test P-Value : 0.8137

```

```

      Sensitivity : 0.9424
      Specificity : 0.8876
    Pos Pred Value : 0.9291
    Neg Pred Value : 0.9080
      Prevalence : 0.6096
    Detection Rate : 0.5746
    Detection Prevalence : 0.6184
    Balanced Accuracy : 0.9150

```

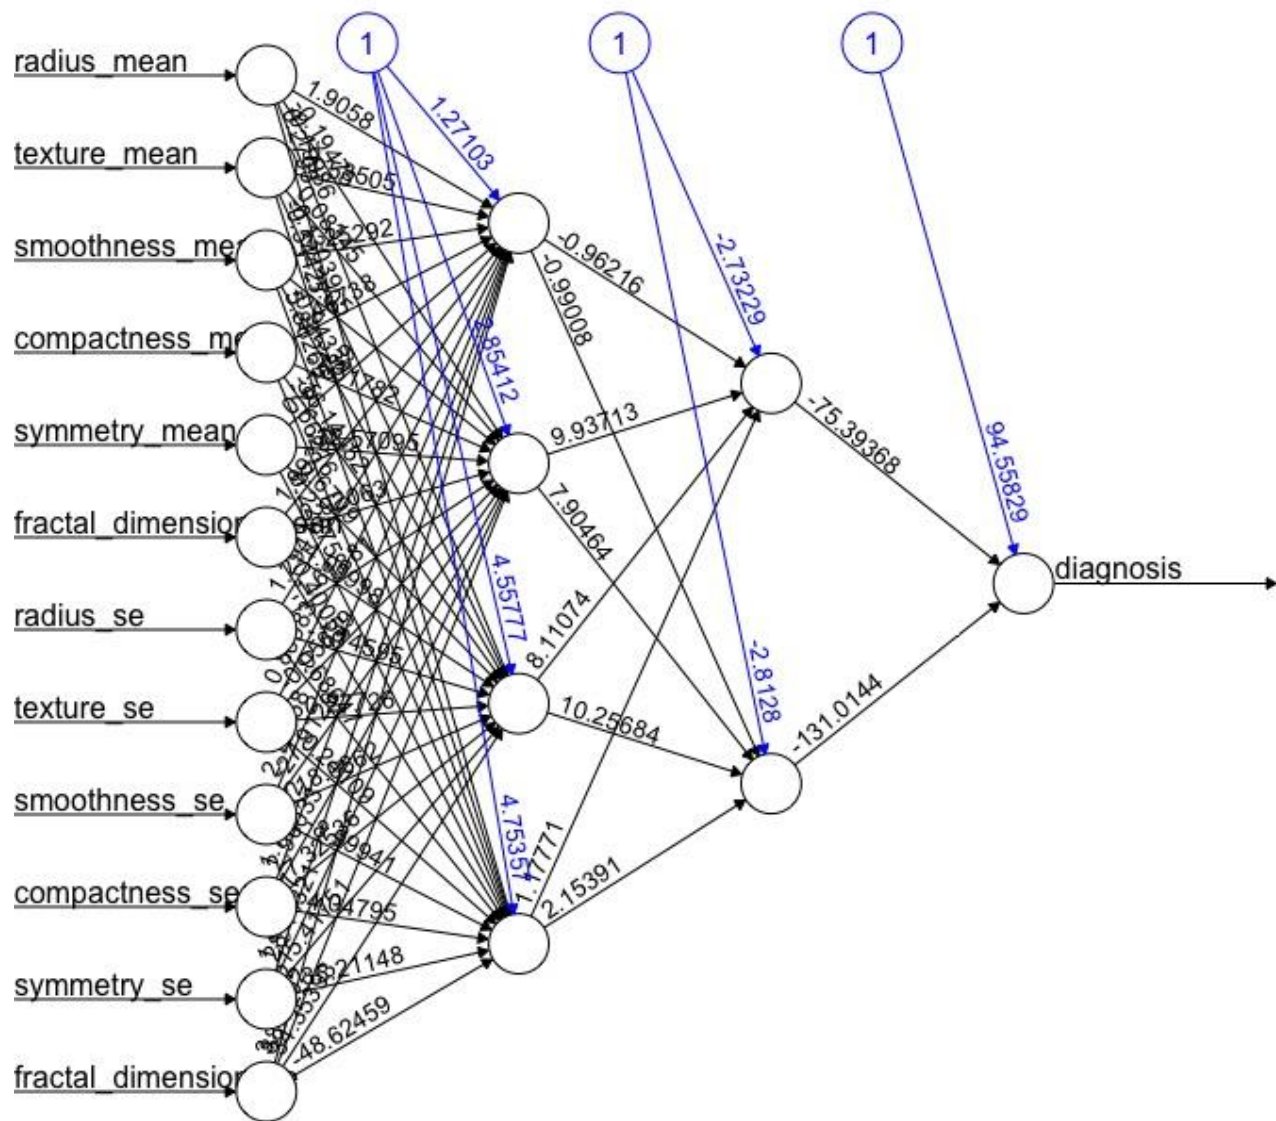
```

'Positive' Class : 0

```

> |

Neural Net	92.11
------------	-------



Results and discussion:

Model	Accuracy %
Logistic Regression	93.42
Random Forest	90.79
KNN	98.19
Neural Network	92.11

It is evident from the table above that KNN model gave the best accuracy in our case. It has the accuracy of 98.19% with significantly lower false positive (1) and false negative (2) as compared to other models.

Note- Visualization and results are attached in the models above.

The output generated by the predictive models will be of great use for the doctors and researchers for finding the right prognosis for Breast Cancer Research.

The final output will be of valuable asset to the researchers as they will be able to predict in advance how soon a benign tumor can convert into a malignant one.

The researchers will be able to have a clear look at the parameters most responsible or the variable contributing most to the tumors for both malign and benign and will be able to make better predictions.

Summary

- 1) Practical implementation of the four models used.
- 2) Even though it is said that the deep learning techniques such as neural network provides the best results, it did not reflect the best result in our case. So, we can conclude that we cannot predict beforehand which model will best work for any dataset until and unless we have run all the different kind of models and evaluate their predicted accuracy.

References:

1. <https://www.kaggle.com/uciml/breast-cancer-wisconsin-data>
2. <https://www.kaggle.com/leemun1/predicting-breast-cancer-logistic-regression>
3. <https://www.hackerearth.com/practice/machine-learning/machine-learning-algorithms/logistic-regression-analysis-r/tutorial/>
4. Class materials